

Parallel Skeletons for Variable-Length Lists in SkeTo Skeleton Library

Aug. 27, 2009

Haruto Tanno

Hideya Iwasaki

The University of Electro-Communications (Japan)

Outline

- Introduction
- Problems of Existing Fixed-Length Lists
- Proposed Variable-Length Lists
 - Skeletons and Operations
 - Data Structure
- Experiments
- Related Work
- Conclusion

Introduction

- Writing efficient parallel programs is difficult
 - synchronization, inter-process communications, data distributions among processes
- Solution: Skeletal parallel programming
 - implements generic patterns within parallel programs
 - C++ parallel skeleton library SkeTo [06]
(intended for distributed environments such as PC cluster)
 - enables users to write parallel programs as if they were sequential


SkeTo

- SkeTo (Skeletons in Tokyo)
 - Constructive parallel skeleton library (C++ and MPI)
 - Joint project of The University of Tokyo, National Institute of Informatics, and The University of Electro-communications
 - Started from 2003, Version 1.0 coming soon
- Distinguishing features of SkeTo
 - It is based on the theory of Constructive Algorithmics
 - It provides skeletons for **lists**, matrices, and trees
 - It introduces no special extension to the base C++


Examples of List Parallel Skeletons

- map skeleton on list
 - applies a function to all elements
- reduce skeleton on list
 - collapses a list with an associative binary operator

map sqr $[1,2,3,4,5,6,7,8]$
 $= [1,4,9,16,25,36,49,64]$



reduce $(+)$ $[1,2,3,4,5,6,7,8]$
 $= 1+2+3+4+5+6+7+8 = 36$



Problem in Existing List

- Computing twin primes

(Pairs of prime numbers that differ by two)

```
nums = [2,3,...,13,14,15], ps = [ ];  
do { // the Sieve of Eratosthenes  
  p = take the front element from nums  
  remove every element that is divisible by p from nums  
  add p to ps  
} while( p <= sqrt_size );  
concatenate ps and nums  
// ps is [2,3,5,7,11,13]  
twin_ps = make pairs of adjacent prime numbers  
remove every pair of prime numbers whose difference isn't two  
// twin_ps = [(3,5), (5,7), (11,13)]
```

List shrinks

List stretches

We can't **shrink** or **stretch** lists because their size is fixed

Problems that Need Variable-Length Lists

Type1: Problems that leave such elements in a given list that satisfy various conditions

- e.g., twin-primes problem, problem of the convex hull

Type2: Searching problems in which the number of candidates for solutions may dynamically change

- e.g., knight's tour

Type3: Iterative calculations in which computational loads for all elements in a list lack uniformity

- e.g., calculations of Mandelbrot and Julia sets

If we can remove elements that have already finished their calculations, we can solve these problems efficiently

Our Proposal

Proposal

- We propose parallel skeletons for lists of variable lengths that enable us to solve a wide range of problems

Approach

- We implement a new library of variable-length lists in SkeTo (They are compatible with existing lists)
 - Provide new **skeletons and operations** that dynamically and destructively change a list's length
 - Change **data structure** that expresses lists
 - Add **automatic data relocation mechanism** for adequate load balancing

Proposed Skeletons and Operations

- **concatmap** applies a function to every element and concatenates the resulting lists
e.g., `concatmap dup [1,2] => [1,1,2,2]`
- **filter** leaves elements that satisfy a Boolean function
e.g., `filter odd [1,2,3,4,5,6,7,8] => [1,3,5,7]`
- **append** concatenates two lists
e.g., `append [1,2] [3,4,5,6] => [1,2,3,4,5,6]`
- **popfront** and **popback** remove an element from the front and the back in a list
- **pushfront** and **pushback** add an element to the front or the back in a list

C++ Program for Twin Prime Problem

```
...  
do{ // the Sieve of Eratosthenes  
    p = nums->pop_front();  
    list_skeletons::filter_ow(IsNotMultipleOf(p), nums);  
    ps->push_back(p);  
} while(p <= sqrt_size);  
ps->append(nums);  
// ps is list of prime numbers  
dist_list<int>* dup_ps = ps->clone<int>();  
dup_ps->pop_front();  
twin_ps = list_skeletons::zip(ps, dup_ps);  
list_skeletons::filter_ow(twin_ps, IsTwin());  
// twin_ps is solution list
```

C++ Program for Knight's Tour

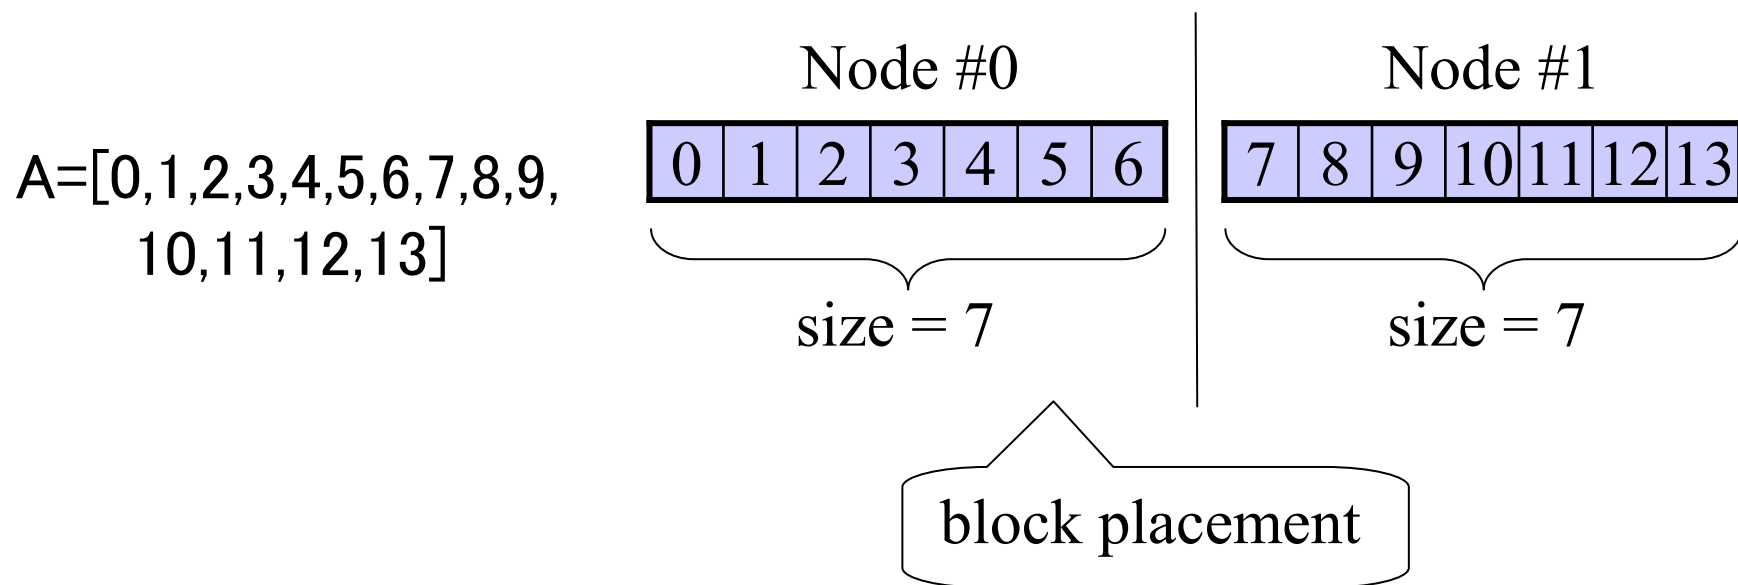
```
dist_list<Board>* bs; // list of solution boards
...
// add initial board state to bs
bs->push_back(initBoard);
// increase bs dynamically up to MAXSIZE
while( bs->get_size() < MAXSIZE ){
    // generate next moves from each board states
    concatmap_ow(nextBoard, bs);
}
// search all solutions with depth first order
concatmap_ow(solveBoard, bs);
// bs is a list of solution boards
```

C++ program for Mandelbrot set

```
dist_list<point>* ps;      // list of points
dist_list<point>* rs;      // list of calculation results
...
for ( int i=0; i<maxForCount; i++ ){
    // progress calculations in small amounts
    map_ow(calc, ps);
    // remove elements that have already finished calculation
    dist_list<point>* es = filtersplit_ow(isEnd, ps);
    // add them to rs
    rs->append(es);
    delete es;
}
rs->append(ps);
// rs is a list of calculation results
```

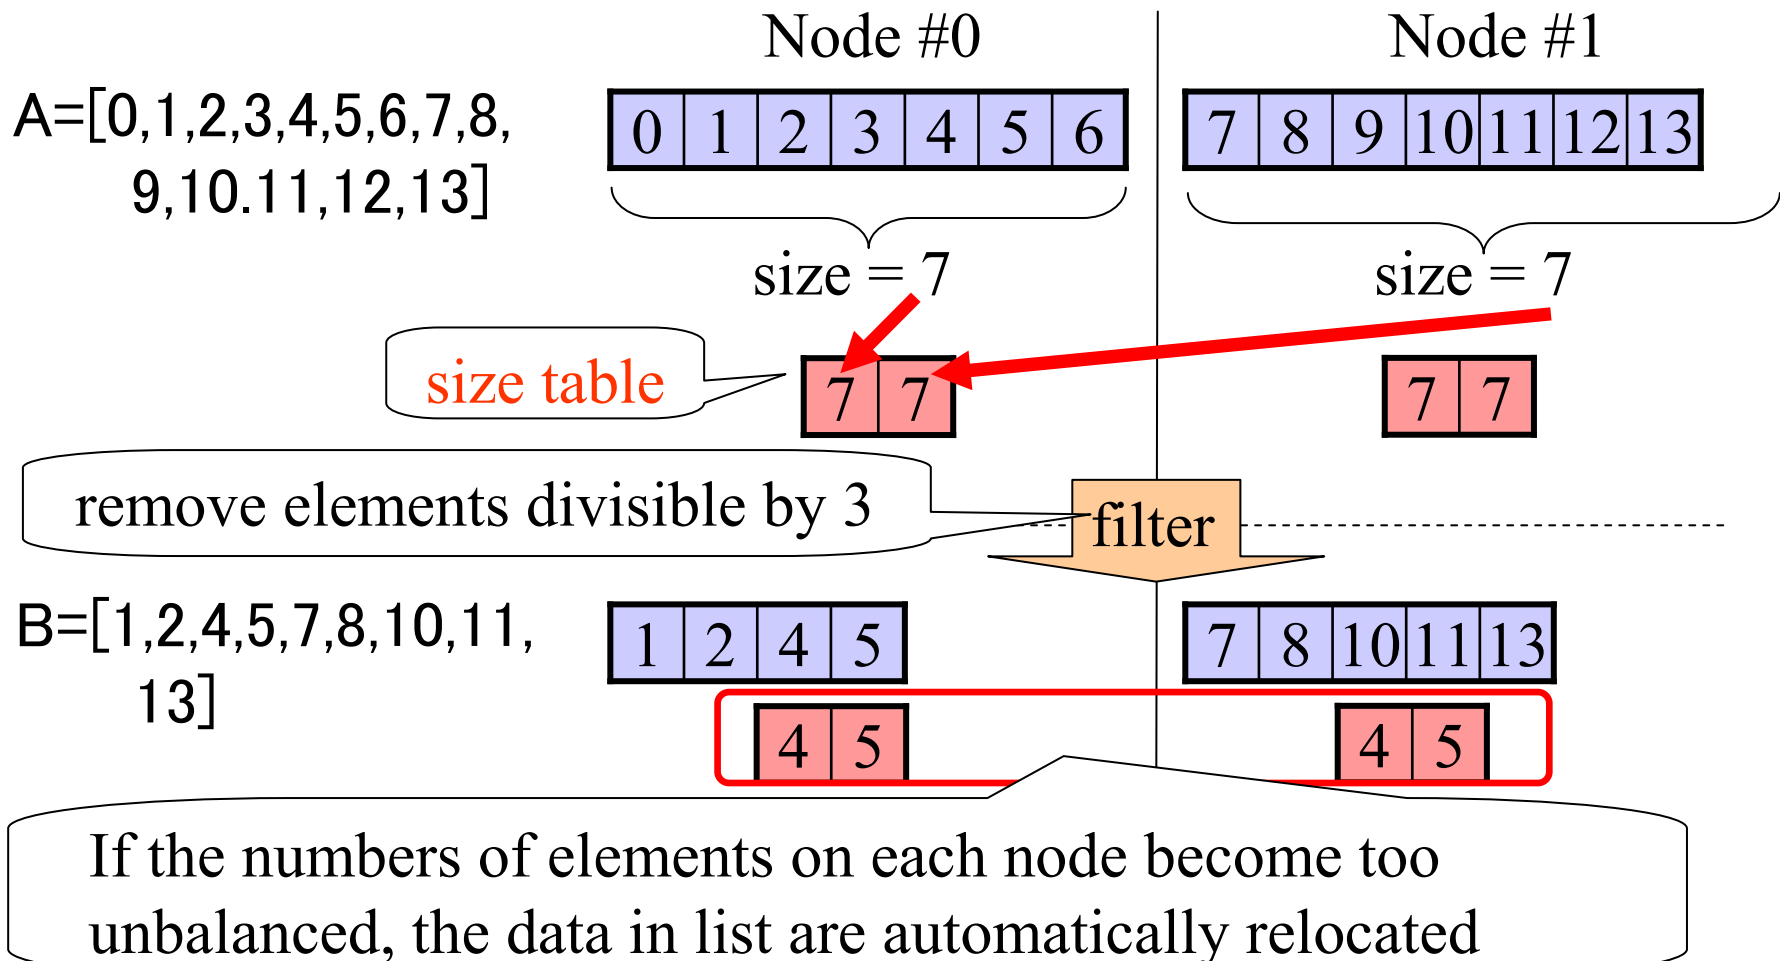
Data Structure of Existing Fixed-Length List

- Elements in lists are equally distributed to each node using block placement
- Data placement does not change during computation



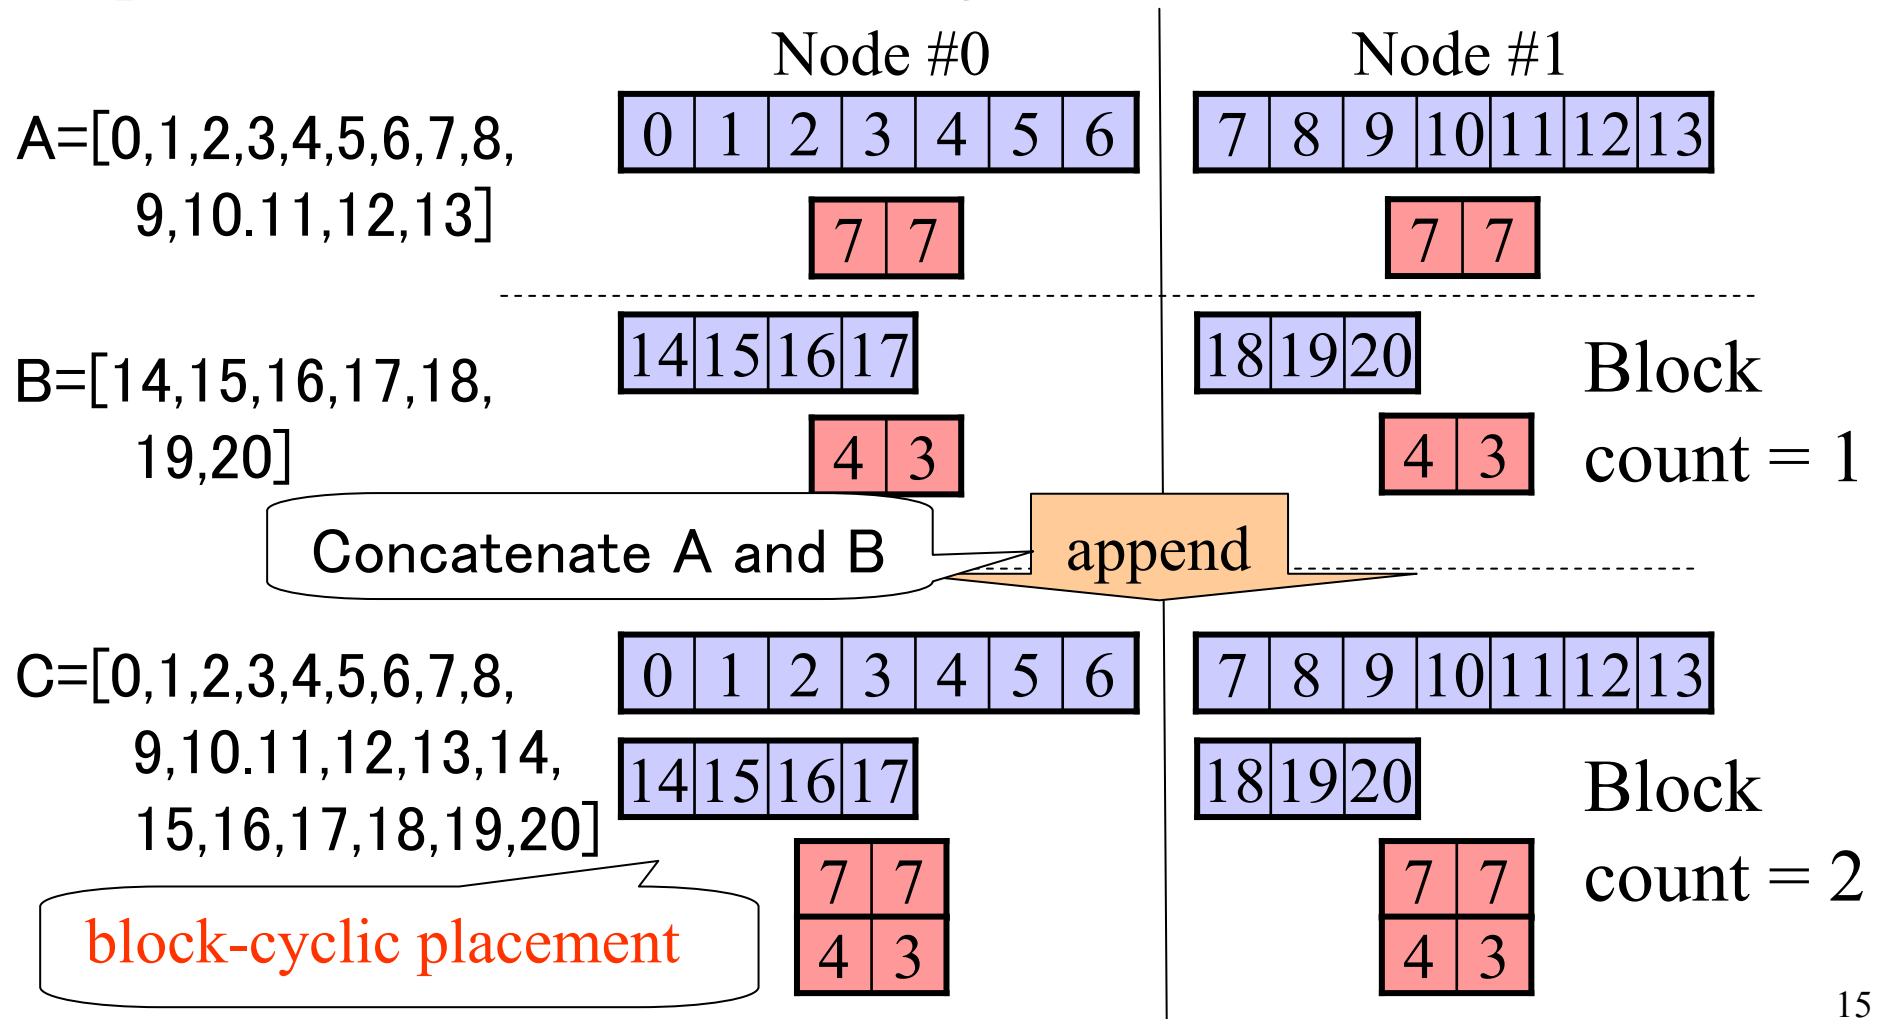
Data Structure of Variable-Length List (1)

- Each node has to know the latest information on the numbers of elements in other nodes



Data Structure of Variable-Length List (2)

- When we concatenate two lists, we adopt block-cyclic placement without relocating the entire amount of data



Micro-benchmark (1)

- Experimental environment
 - Pentium4 3.0GHz, Mem 1GB, Linux 2.6.8, 1000BaseT Ethernet
- We measured the execution times for applying **map**, **reduce**, and **scan**, and the data relocation in a list
 - Input list: 80 million elements
 - block count: from 1 to 4,000
 - two functions: short/long execution time

Micro-benchmark (2)

Execution time (s)

block count	1	10	100	1000	2000	3000	4000
Map (short)	0.0128	0.0129	0.0128	103%	0.0129	0.013	0.0132
Reduce (short)	0.0183	0.0182	0.0183	109%	0.0194	0.0197	0.0200
Scan (short)	0.0407	0.0408	0.0411	142%	0.0484	0.053	0.0580
Map (long)	16.8	16.8	16.8	100%	16.8	16.8	16.8
Reduce (long)	16.9	16.9	16.9	101%	16.9	16.	17.0
Scan (long)	33.8	33.8	33.8	101%	33.9	34.	34.1
Data Reloaction	-	3.74	4.64	4.67	4.66	4.62	4.72

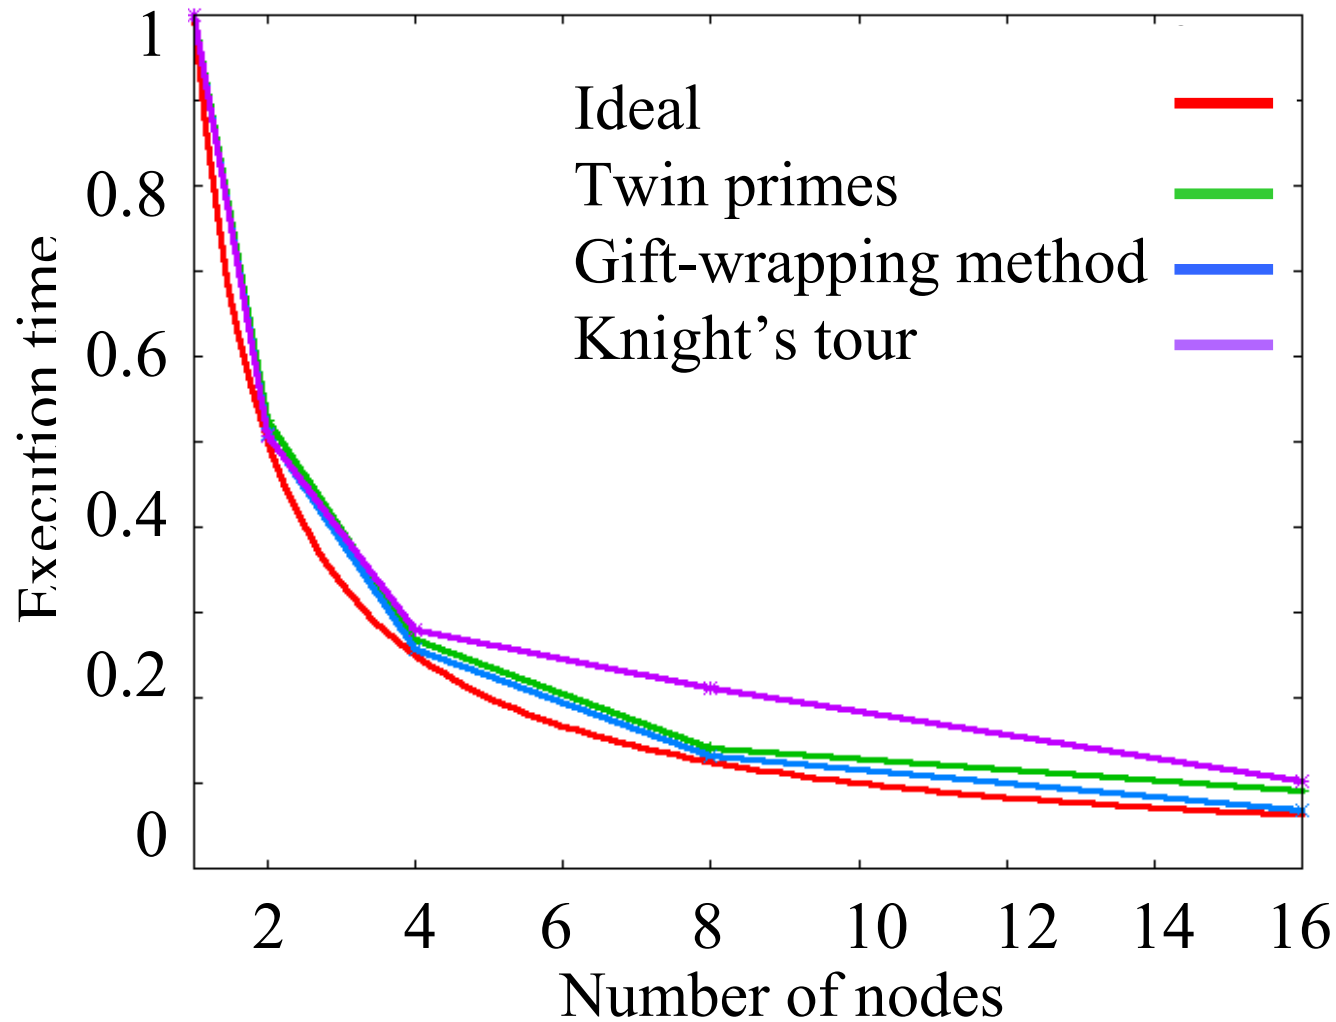
The overheads of data relocation are large

➔ It is effective to delay the relocation of data

Macro-benchmark (1)


- Type1
 - Twin primes (list of 10 million integers)
 - Gift-wrapping method (1 million points)
- Type2
 - Knight's tour (5×6 board)
- Type3
 - Mandelbrot set ($1,000 \times 1,000$ coordinates)
 - Using variable-length lists with 100 iterative calculations $\times 100$ times
 - Using fixed-length lists with 10,000 iterations

Macro-benchmark (2)



These results indicate excellent performance in all problems

Macro-benchmark (3)



Programs with variable-length lists show good speedups

Parallel Skeleton Libraries

- P3L [95], Muesli [02], Quaff [06]
 - support data parallel skeletons
 - offer lists (distributed one-dimensional arrays)
- eSkel [05]
 - supports task parallel skeletons but does not support data parallel skeletons for list like map and reduce
- Muskel [07], Calcium [07]
 - a Java skeleton library on a grid environment

These libraries do not support variable-length lists

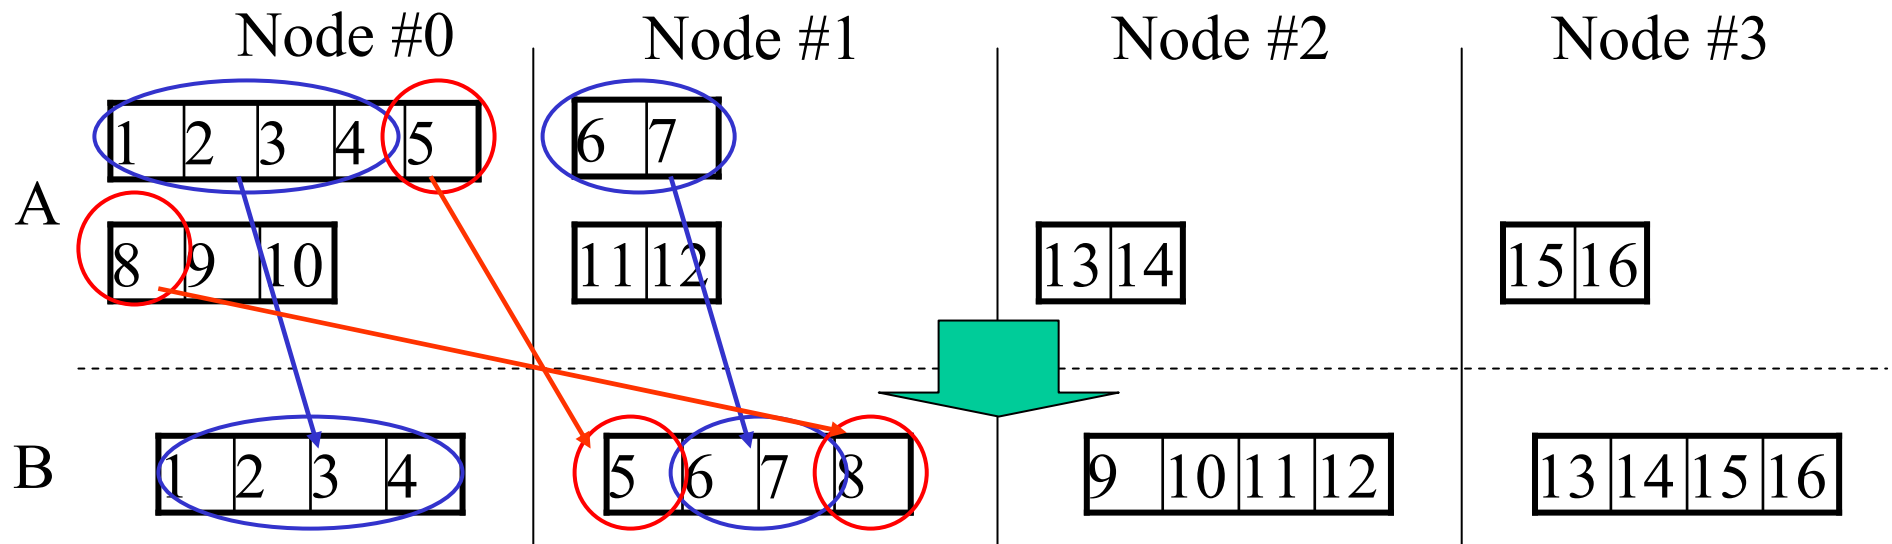
Another Group of Libraries

- STAPL [07]
 - Offers variable-length arrays and lists (pVector, pList)
 - Provides the same operations as C++ STL
 - Does not have operations such as concatmap
- Data Parallel Haskell [08]
 - Offers distributed nested lists
 - Provides filter, concatmap, and append
 - Only targets shared-memory environments

Conclusion

- We proposed parallel skeletons for variable-length list and their implementation
 - We proposed skeletons and operations for variable-length list, e.g., `concatmap`, `filter`, and `append`
 - We adopted a block-cyclic representation of lists with size tables
 - We confirmed the efficiency of our implementation through tests in various experiments

Data relocation (1)

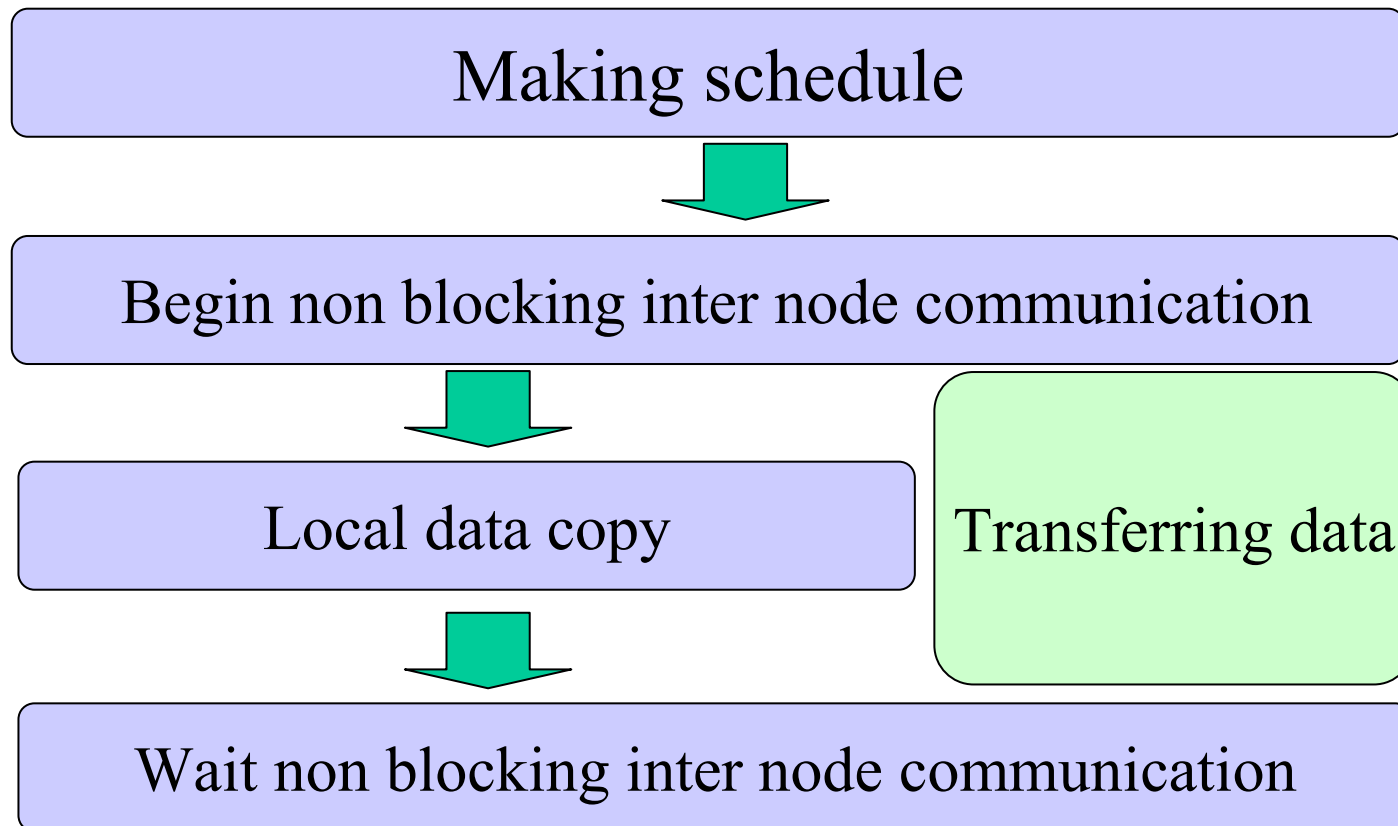


source	dest	
#0	#0	(1,2,3,4)
#0	#1	(5)
#1	#1	(6,7)
#0	#1	(8)
...

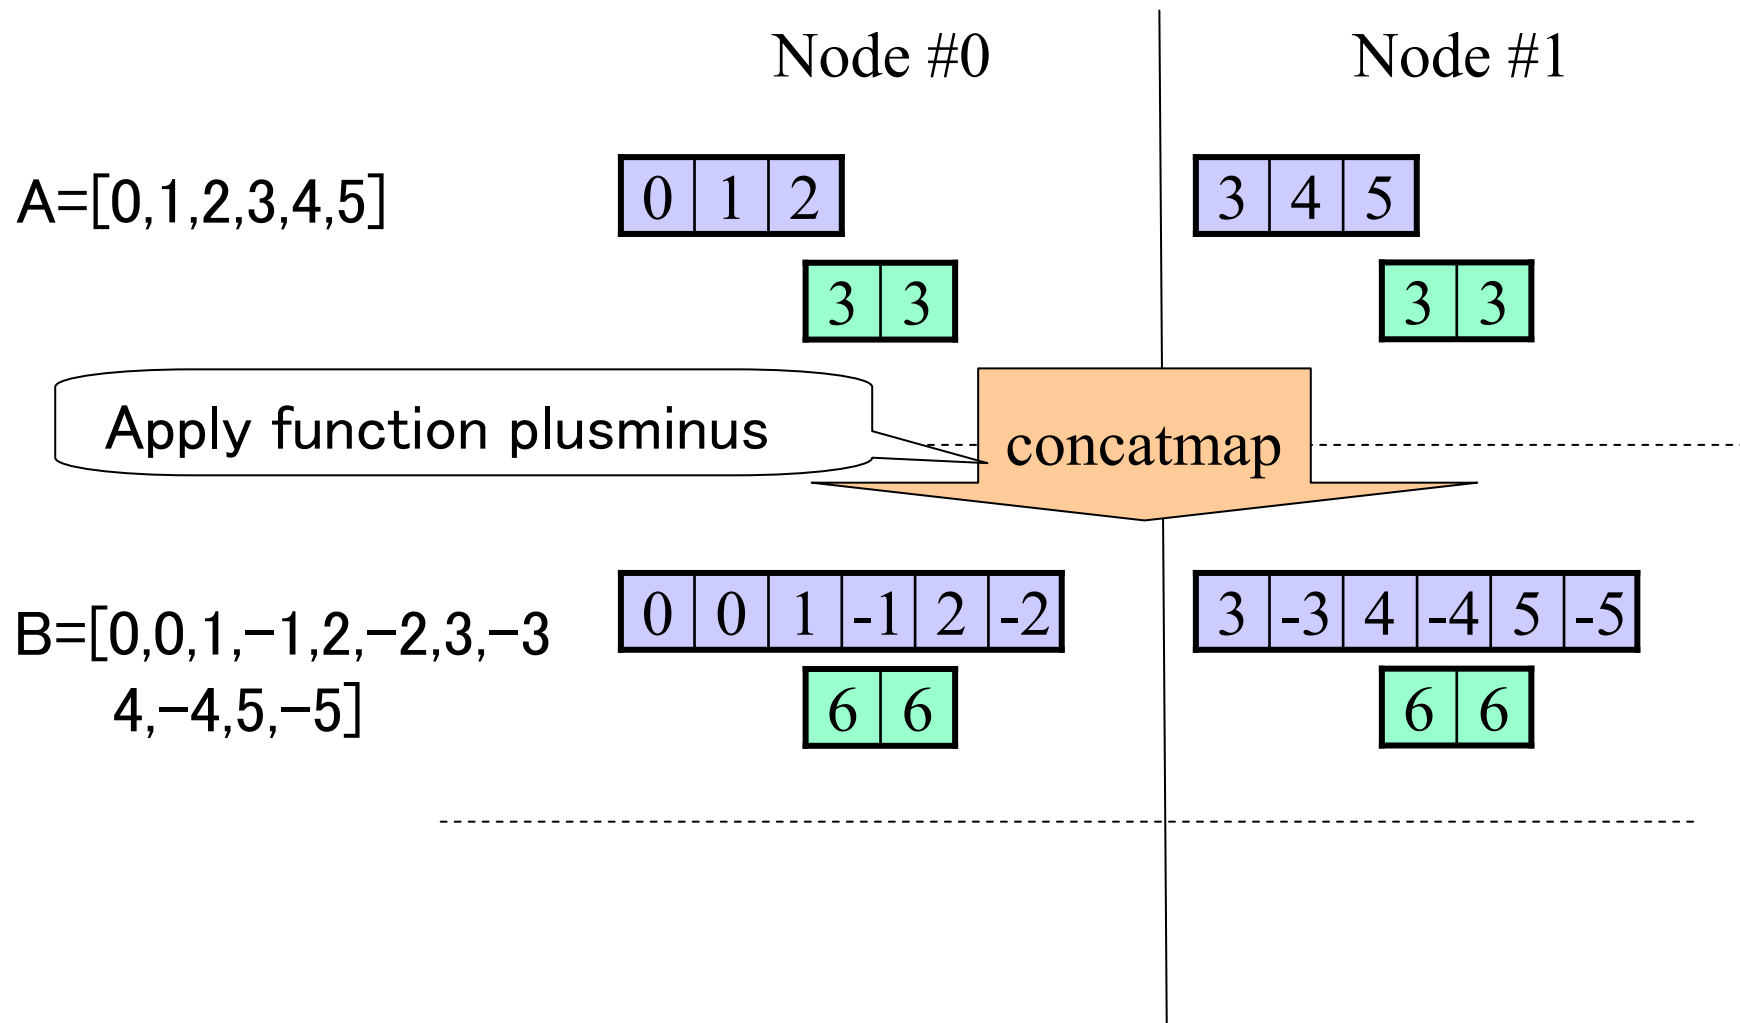
local data copy

inter node communication

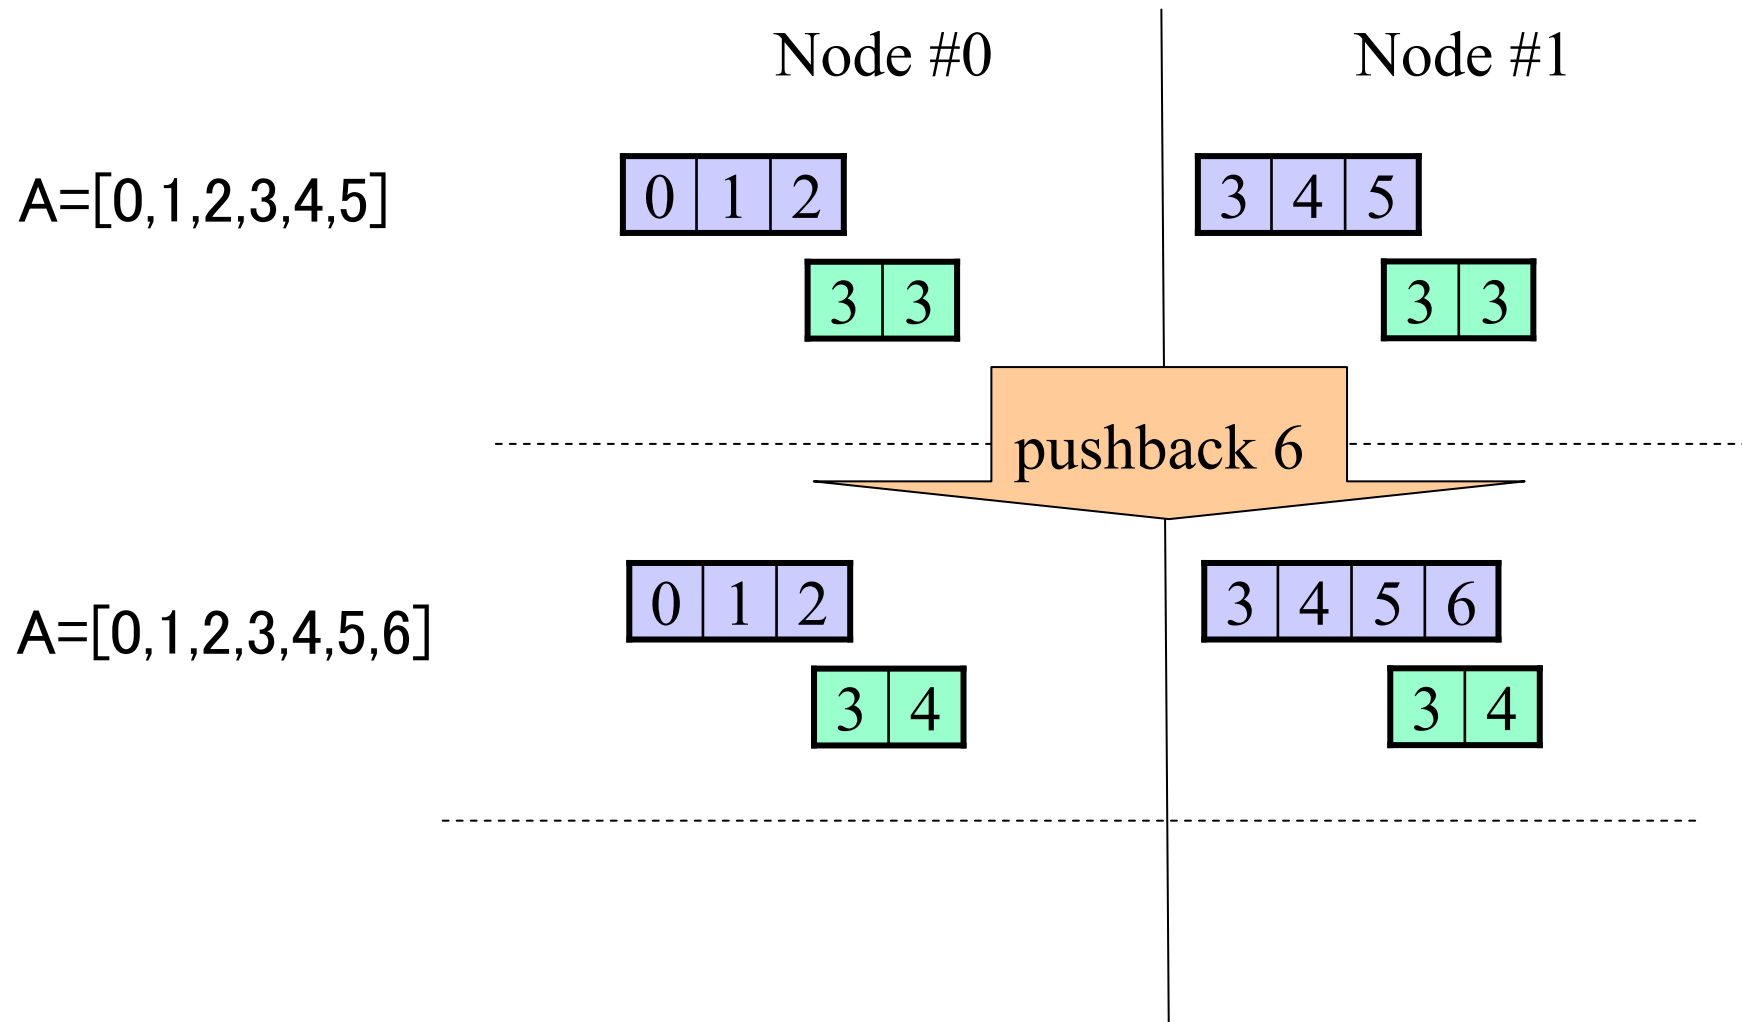
Data relocation (2)



Cocatmap



Pushback



Condition of Data Relocation

$$\text{imbalance}(A) = n \times \max(P_1, \dots, P_n) / (P_1 + \dots + P_n)$$


where $P_i = \sum_{j=1}^m p_{ij}$, n = the number of nodes, m = the block count of A ,
 p_{ij} = the number of elements of the j -th block at the i -th node.

$$\text{imbalance}(A) > 1.5$$

➔ data relocation

Convex hull

- Gift wrapping method



Scan in Variable-Length List

Node #0	Node #1	Node #2	Node #3
<div>0 0 0</div> <div>3 3 3</div>	<div>0 1 1</div> <div>3 4 4</div>	<div>1 1 2</div> <div>4 4</div>	<div>2 2 2</div> <div>5 5</div>

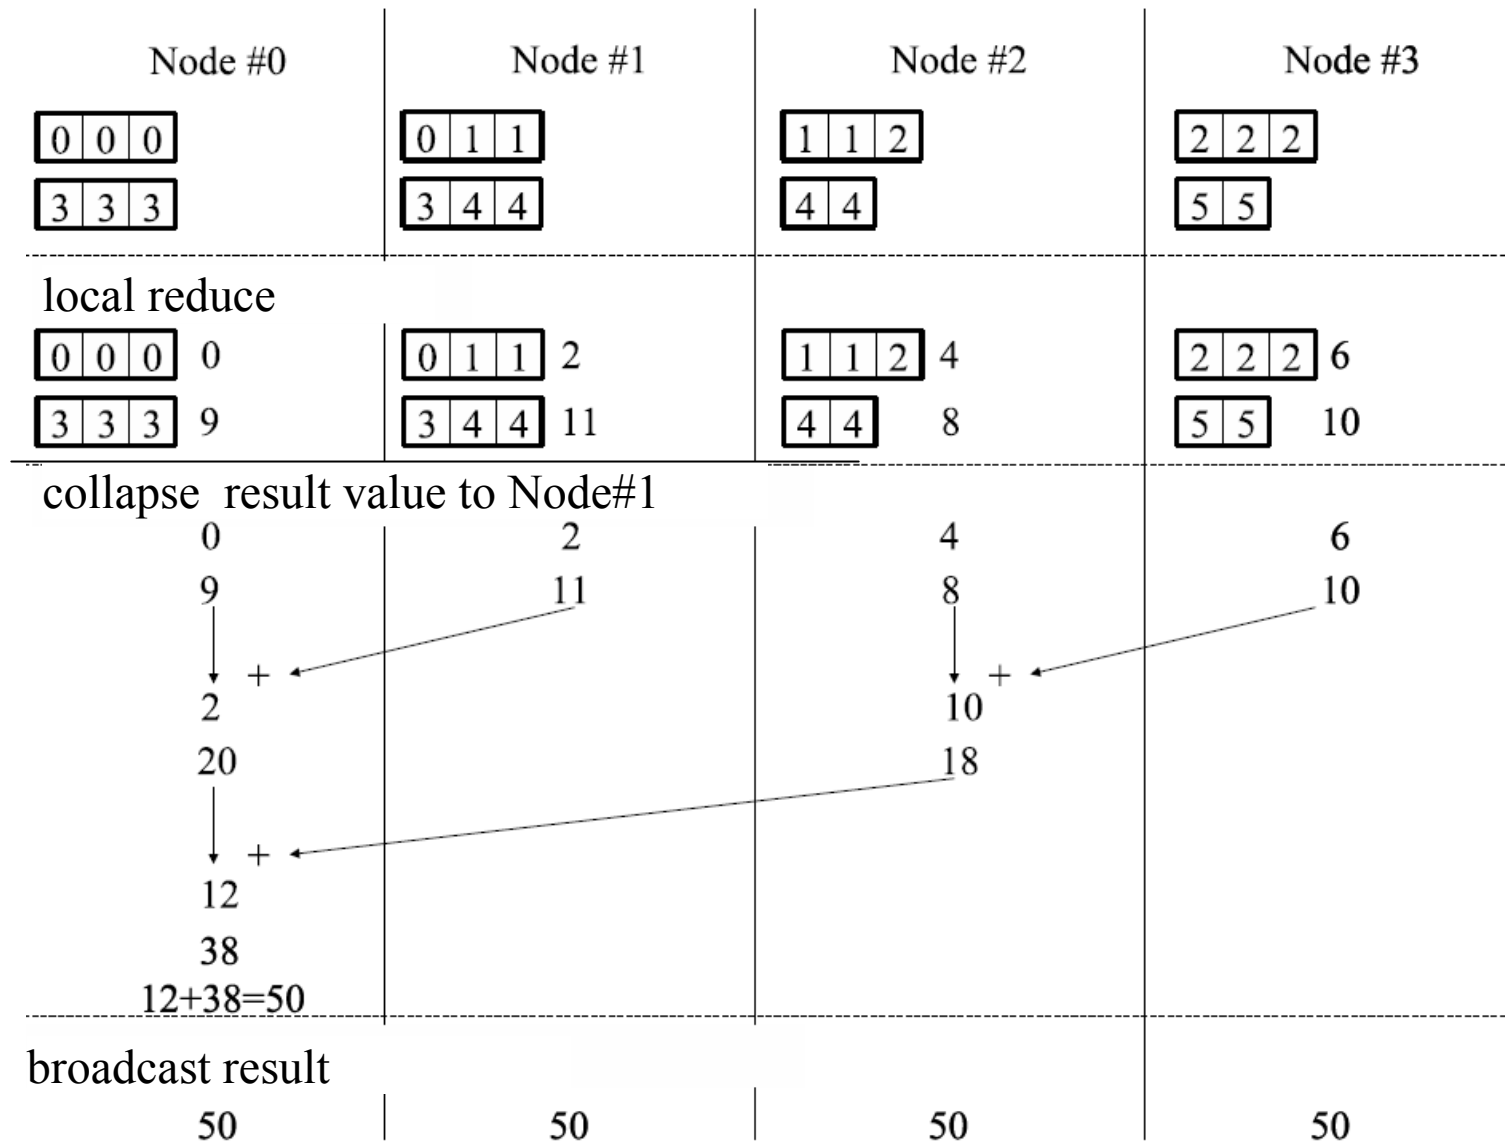
local reduce			
<div>0 0 0</div> 0 <div>3 6 9</div> 9	<div>0 1 2</div> 2 <div>3 7 11</div> 11	<div>1 2 4</div> 4 <div>4 8</div> 8	<div>2 4 6</div> 6 <div>5 10</div> 10

sharing local reduce among nodes			
<div>0 0 0</div> 0 2 4 6 <div>3 3 3</div> 9 11 8 10	<div>0 1 2</div> 0 2 4 6 <div>3 7 11</div> 9 11 8 10	<div>1 2 4</div> 0 2 4 6 <div>4 8</div> 9 11 8 10	<div>2 4 6</div> 0 2 4 6 <div>5 10</div> 9 11 8 10

local scan			
<div>0 0 0</div> - 0 2 6 <div>3 3 3</div> 12 21 32 40	<div>0 1 2</div> - 0 2 6 <div>3 7 11</div> 12 21 32 40	<div>1 2 4</div> - 0 2 6 <div>4 8</div> 12 21 32 40	<div>2 4 6</div> - 0 2 6 <div>5 10</div> 12 21 32 40

calculating final scan			
<div>0 0 0</div> <div>15 15 15</div> add 12	<div>0 1 2</div> add 0 <div>24 28 32</div> add 21	<div>3 4 6</div> add 2 <div>36 40</div> add 32	<div>8 10 12</div> add 6 <div>45 50</div> add 40

Reduce in Variable-Length Lists



Zip in Variable-Length Lists

